

Report Phase 2:- Outlaws [Kartik Garg (2019101060)]

Optimization Algorithms/Heuristics Used

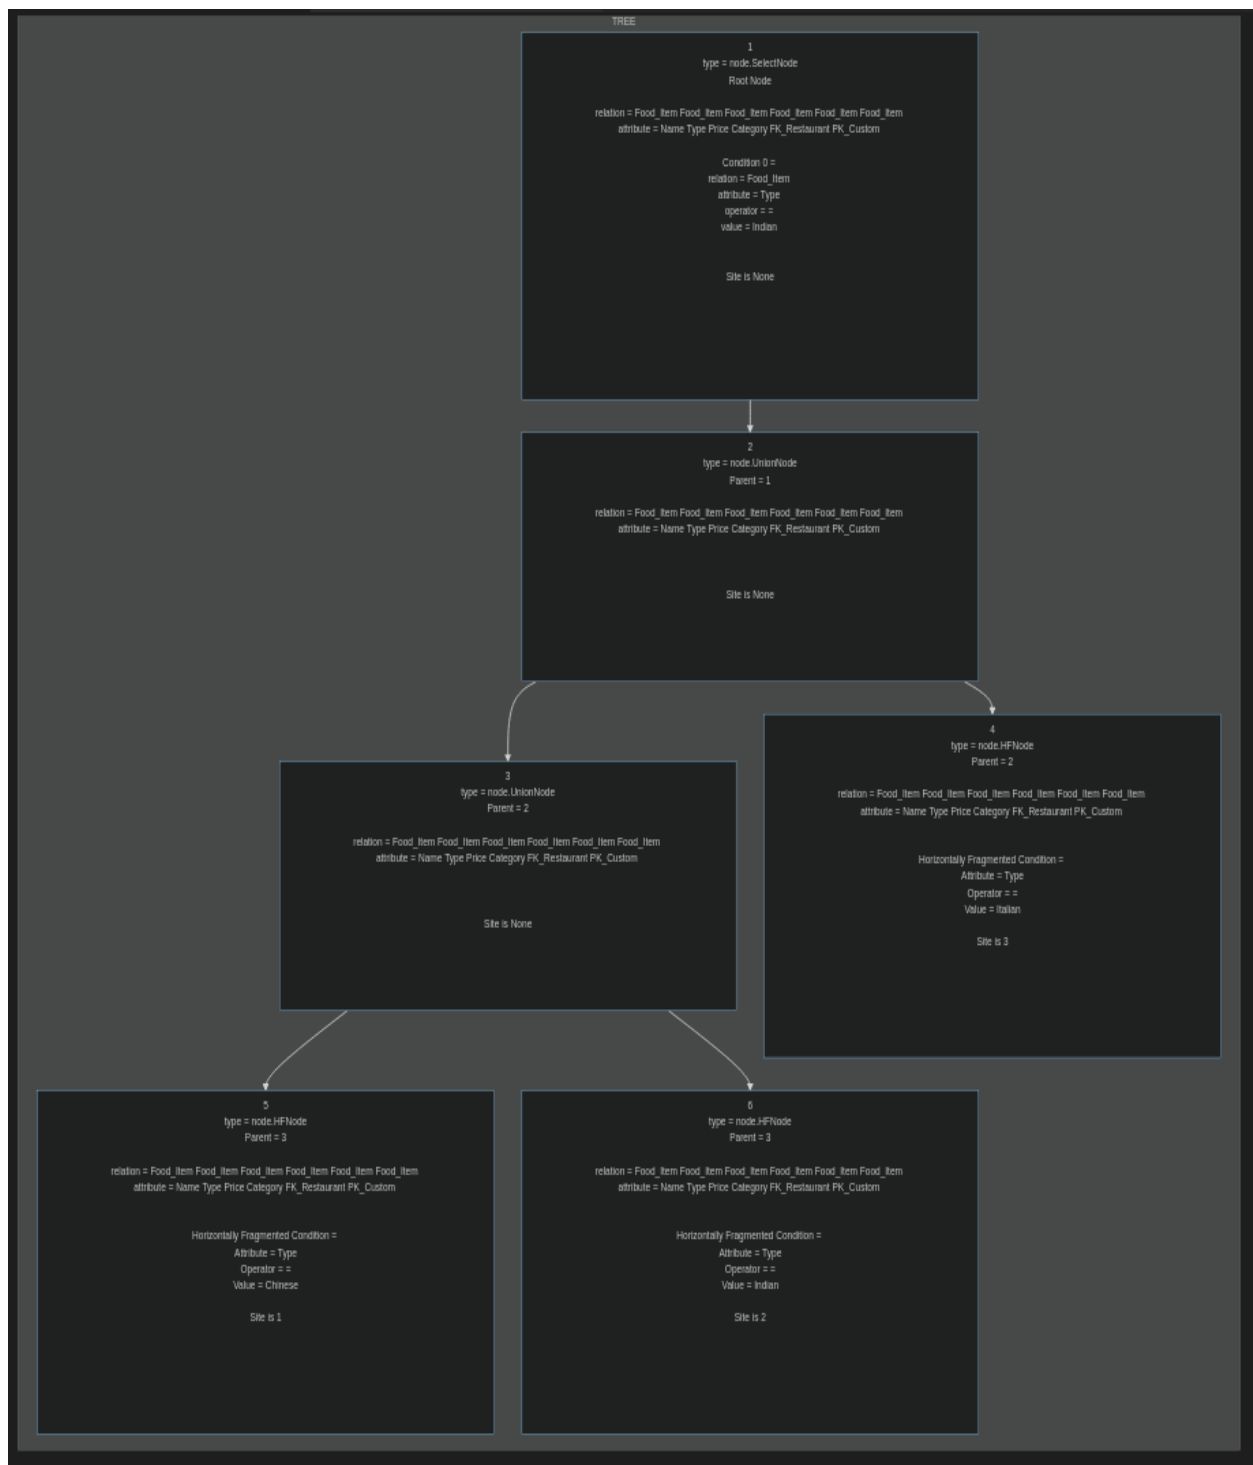
Heuristics

Standard heuristics like pushing the projects and selects down the tree are already present. And not just pushing down for each select condition we also check whether the relation/fragment node on which that select is going to sit on top can logically have the answer or not. So I have a method implemented in my Query class called checkCondition which basically takes two conditions and compares them and outputs whether they can have tuples which are common to both. So if we encounter a condition which is logically contradicting the one which we have for, lets say, Horizontal Fragment Node. In this case we will eliminate this fragment straight away as it cannot have any result tuples inside it. So, this kind of optimization is there.



Above attached is the original graph which is made for the query
Select * from Food_Item where (Type='Indian')

When we break relations into fragments it becomes like the one whose image is on the next page.



But after passing the select condition down, it becomes a lot more simpler graph, image attached on next page

TREE

1

type = node.SelectNode

Root Node

relation = Food_Item Food_Item Food_Item Food_Item Food_Item Food_Item

attribute = Name Type Price Category FK_Restaurant PK_Custom

Condition 0 =

relation = Food_Item

attribute = Type

operator = =

value = Indian

Site is None



2

type = node.HFNode

Parent = 1

relation = Food_Item Food_Item Food_Item Food_Item Food_Item Food_Item

attribute = Name Type Price Category FK_Restaurant PK_Custom

Horizontally Fragmented Condition =

Attribute = Type

Operator = =

Value = Indian

Site is 2

Because we were able to identify from the conditions and decide which fragment will have the final result, we were able to transform a much complex query tree to the one which is much simpler and contains only two nodes.

Above query is executed and gives the following output

```
->> select * from Food_Item where (Type='Indian');
```

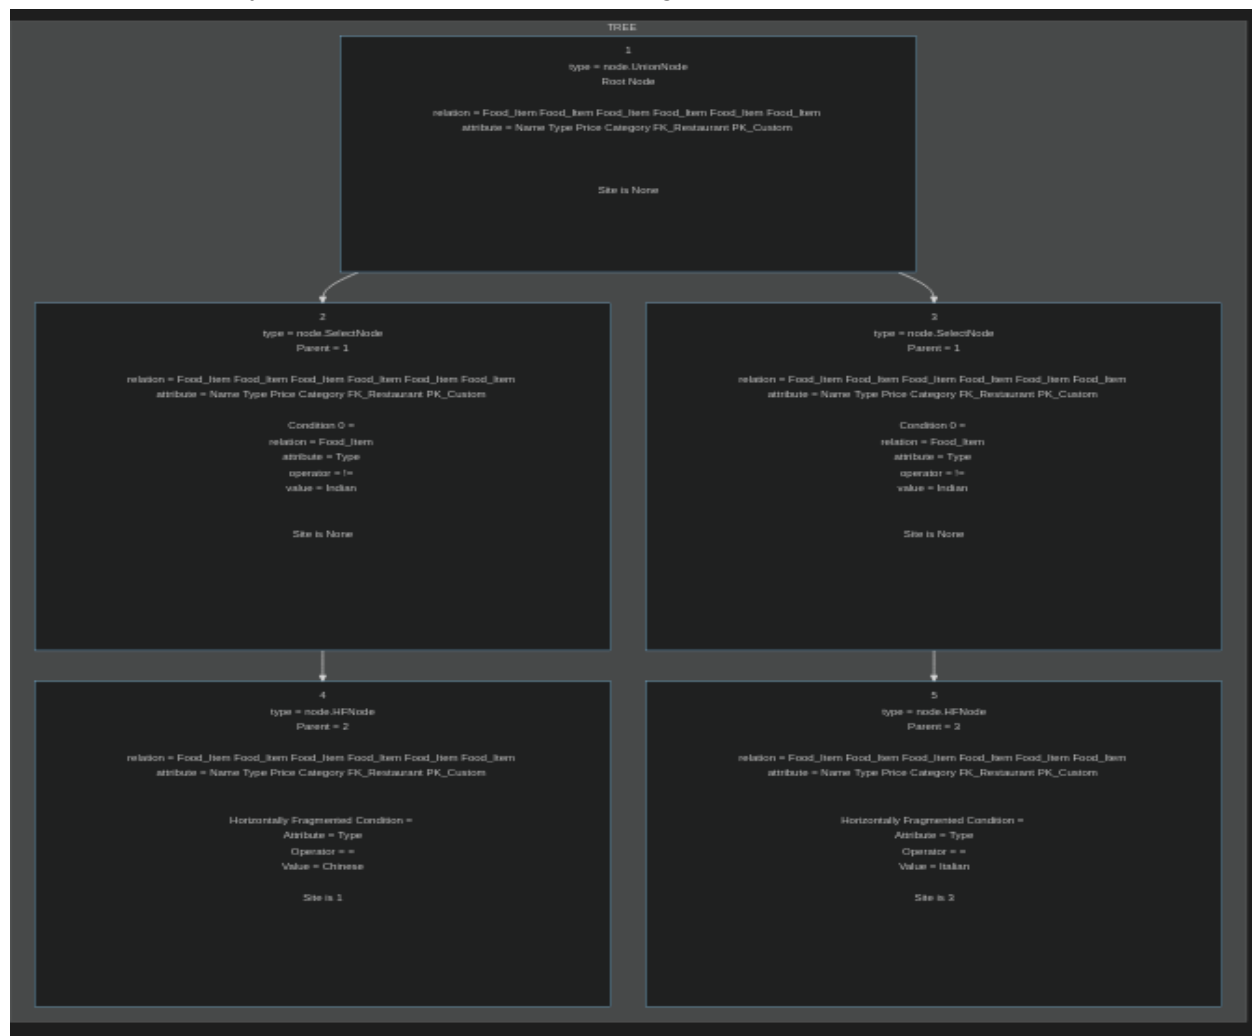
	Name	Type	Price	Category	FK_Restaurant	PK_Custom
0	Special_Shawarma	Indian	100	Non-Veg	AAI	AAP
1	Cheese_Shawarma	Indian	90	Veg	AAI	AAQ
2	Honey_Chilli	Indian	120	Veg	AAK	AAV

Time Taken (In Seconds): 1.8483760356903076

If the query is changed to a query

Select * from Food_Item where (Type!='Indian')

Now, the final query tree contains the other two fragments, like shown below



And the Output goes like this,

```
->> select * from Food_Item where (Type!='Indian');
```

	Name	Type	Price	Category	FK_Restaurant	PK_Custom
0	Momos	Chinese	40	Veg	AAJ	AAL
1	Manchurian	Chinese	50	Veg	AAK	AAM
2	Noodles	Chinese	45	Veg	AAK	AAN
3	Chowmein	Chinese	70	Veg	AAK	AAO
4	Marghrita	Italian	100	Veg	AAH	AAR
5	Non_Veg_Loaded	Italian	150	Non-Veg	AAH	AAS
6	Veg_Loaded	Italian	125	Veg	AAH	AAT
7	Pasta	Italian	95	Veg	AAH	AAU

```
Time_Taken (In Seconds): 4.057796955108643
```

So, this was all about the basic heuristics. Here we pushed the selects and projects down and removed the irrelevant fragments which did not contribute to our final result.

Optimizations

First optimization which we have in our code is that, before making the query graph we find the best join order based on the selectivities provided to us. What we do is, we try all the join orders which are possible for a query and find out the one which has the minimum cost. The code snippet for the same is attached below. Here, we pass a join order to this function and we calculate the cost associated with it. We keep track of which join order gave me the minimum cost and we proceed with selecting that join order.

```
def calcCost(self, joinOrder, relations):  
    '''  
    Given a join order, join selectivities and relation sizes  
    this function is supposed to calculate the cost of their joining  
    '''  
    cost = 0  
    relation_sz = {}  
    for x in relations:  
        relation_sz[x] = config.relationNumEntries[x]  
  
    for x in joinOrder:  
        r1 = x[0]  
        r2 = x[0]  
  
        joinSel = config.joinSelectivities[(r1,r2)]  
        sz1 = relation_sz[r1]  
        sz2 = relation_sz[r2]  
        nusz = sz1*sz2  
  
        cost = cost + nusz  
        relation_sz[r1] = nusz*joinSel  
        relation_sz[r2] = nusz*joinSel  
    return cost
```

Other optimizations which make use of join selectivity factors and the allocation information are deployed while handling joins and unions.

So, in **unions** we check the length of the two relations which are taking part in this union. And we will ship the one with fewer rows to the relation with more rows. This adds up to our **second optimization**.

For shipping the relations a helper utility file is written which has 4 functions in it which are dumpTable, copyFromServer, copyToServer, importTable. The first utility function dumps a table as a text file which contains sql statements to generate the exact same table. The second function brings this txt file to my machine. The third function will copy this .txt file from my machine to the server which will be specified. The last utility function will basically import the table from this txt file into the mysql server.

For **optimizing the joins**, we have made use of semi-joins as specified in the description document. So we have to first find out which relation should be shipped to the site of which relation (given both the relations are in separate sites, if they are on the same site then we just execute the normal inner join). To find out this we have designed some preprocessing functions which extract certain information about all the relations which is necessary to calculate the cost.

```
def dumpTable(relname,siteno,ofile = 'dump.txt'):
    """
    Function which will dump a mysql table using mysqldump command into ofile
    """
    config.logger.log("utility::dumpTable")

    command = "echo 'use "+config.catalogName+";' > ./Outlaws/"+ofile
    (a,b,c) = config.paramikoConnections[siteno].exec_command(command)
    op = c.read()

    command = "mysqldump -u user -piiit123 "+config.catalogName+" "+relname+" >> ./Outlaws/"+ofile
    # print(command)
    (a,b,c) = config.paramikoConnections[siteno].exec_command(command)
    op = c.read()
    # print(op)
```

```
def copyFromServer(siteno,ofile = 'dump.txt'):
    """
    Function to copy a file from the Server to my machine
    """
    config.logger.log("utility::copyFromServer")

    with SCPClient(config.paramikoConnections[siteno].get_transport()) as scp:
        scp.get("./Outlaws/"+ofile,ofile)
```

```
def copyToServer(siteno,ofile = 'dump.txt'):
    """
    Function to copy a file from my machine to Server
    """
    config.logger.log("utility::copyToServer")

    with SCPClient(config.paramikoConnections[siteno].get_transport()) as scp:
        scp.put(ofile,"./Outlaws/"+ofile)
```

```
def importTable(siteno,ofile = 'dump.txt'):
    '''
    Function to read the dump.txt file
    '''
    config.logger.log("utility::importTable")

    command = "mysql -u user -p111t123 < ./Outlaws/"+ofile
    (a,b,c) = config.paramikoConnections[siteno].exec_command(command)
    op = c.read()
```

There are mainly two preprocessing function which are:

1. getEntrySizes()
2. getRelationLengths()

Both these functions communicate with each of the server sites and then using the sql's information_schema we get the relation length, in this case it can also be fragments, and the individual column sizes. From individual column sizes we can calculate the tuple size by adding all the size values for each of the column values

And the relation lengths are directly stored in the information_schema.TABLES relation and are retrieved from there only.

Now when we calculate the cost of A join B as (B semi join A) join A, we do it by using the formula given below

$$\text{Cost} = CP * \text{Len}(B) + CTF * \text{Len}(\text{Project Join Attr}) * \text{Size}(\text{Join Attr B}) + CP * L * \text{Len}(A) + CTF * \text{Join Selectivity}(A,B) * L * \text{Len}(A) * \text{Size}(A)$$

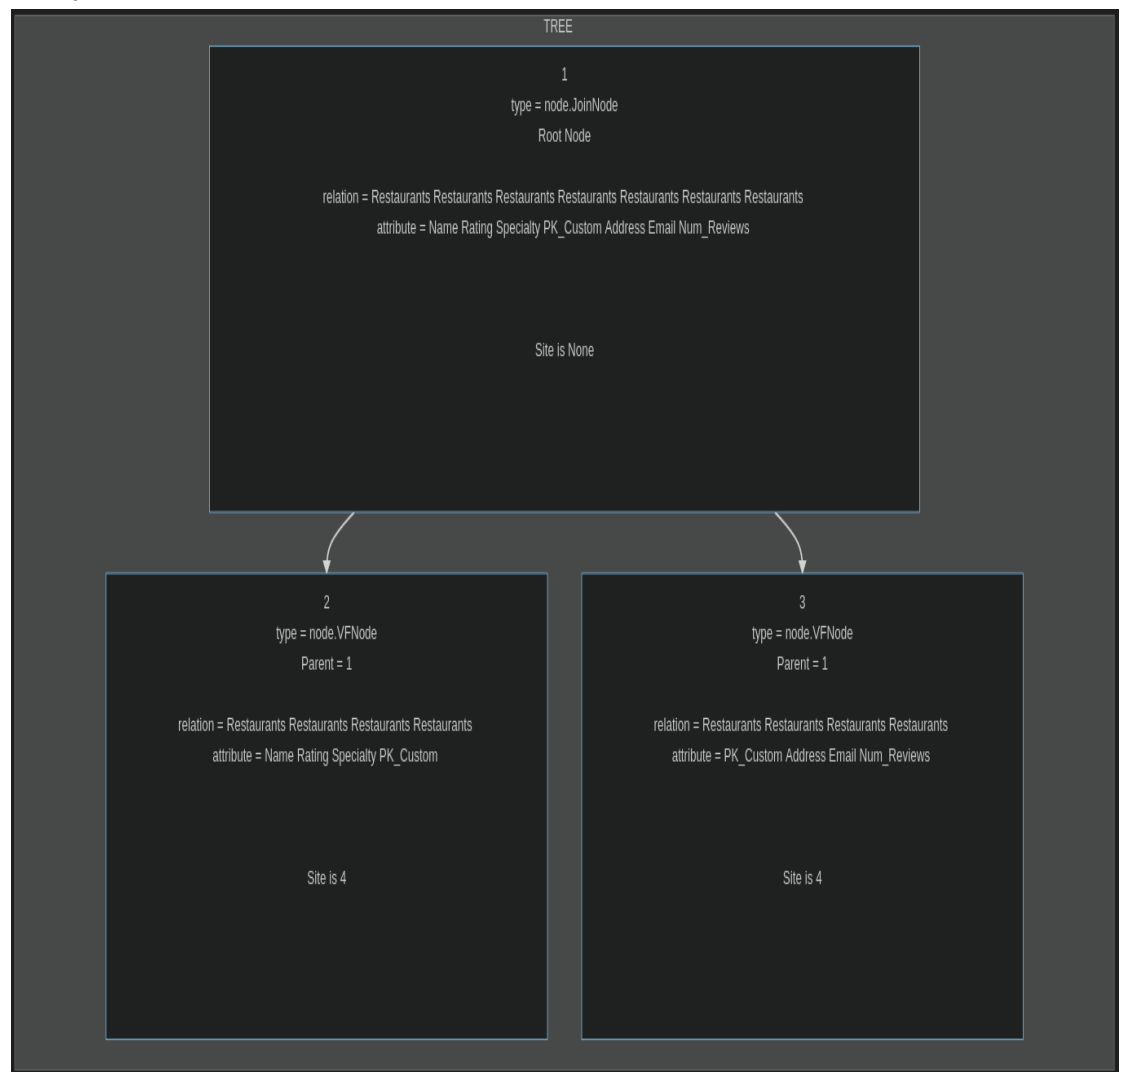
$CP = \text{Processing Cost}$ $CTF = \text{Transfer Cost between the two sites in consideration.}$

In a similar fashion we calculate the (A semi join B) join B variant and we pick the choice which gives us the minimum cost and we execute it. This was all about our **third optimization**.

Optimized Query Tree and Final Output

It is already attached in the above document for some of the queries. I am also attaching it below in this section for a couple of more queries which will involve use of joins and aggregate operators.

1. Select * from Restaurants;
 - a. Query Tree



- b. Result

```
->> select * from Restaurants;
```

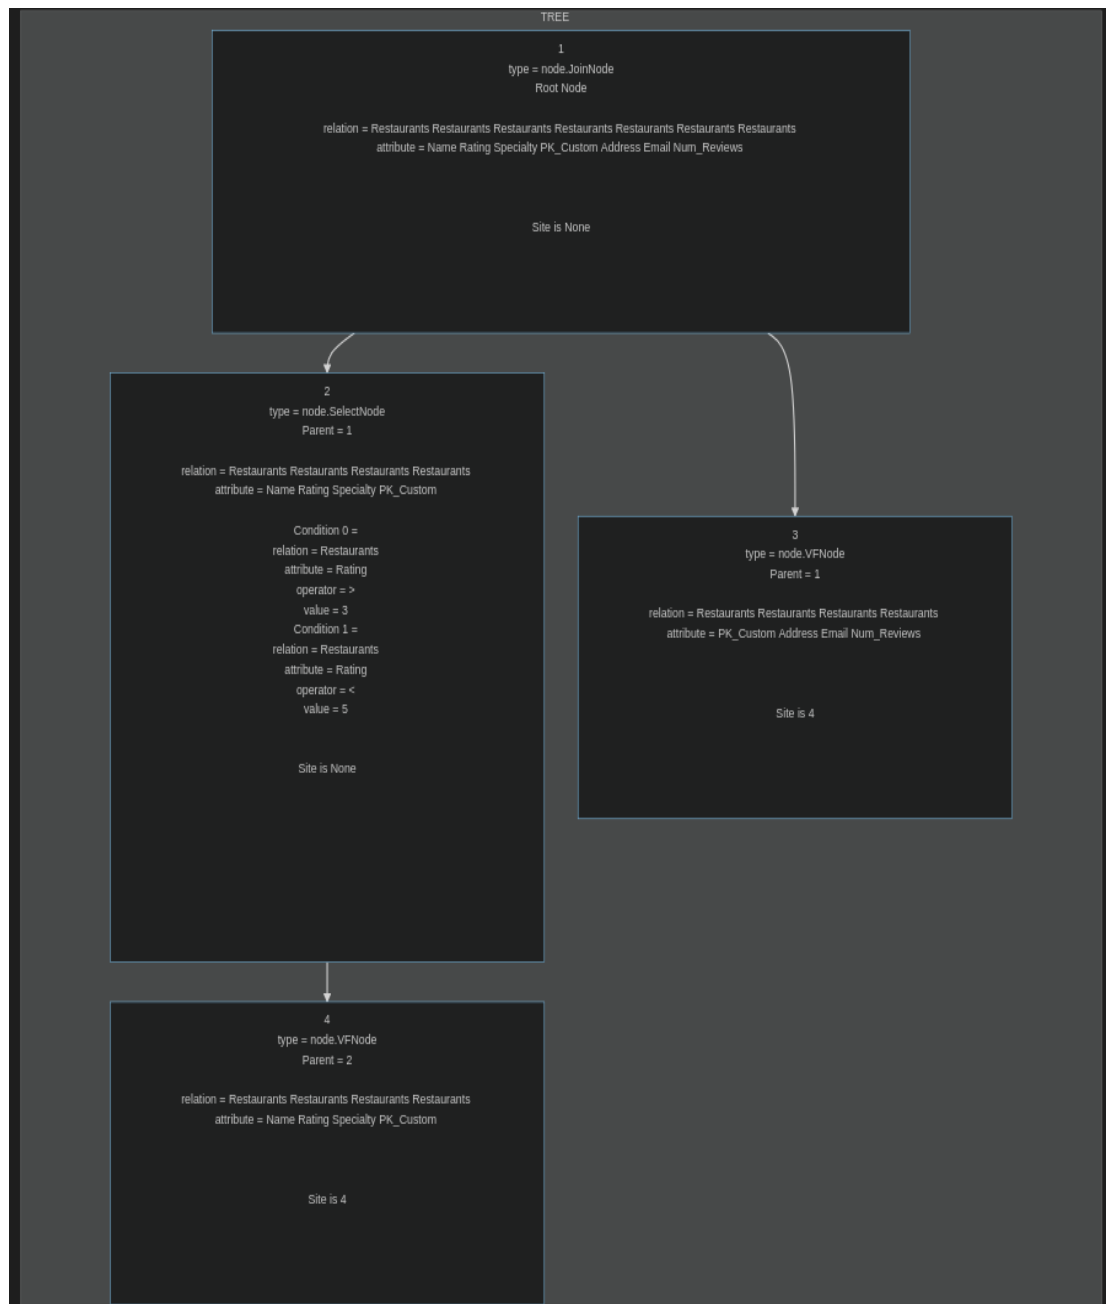
	Name	Rating	Specialty	PK_Custom	Address	Email	Num_Reviews
0	Dominos	4	pizza	AAH	Gachibowli	domi@gmail.com	45
1	SKShawarma	3	shawarma	AAI	DLF	sks@gmail.com	38
2	WOWMomos	5	momos	AAJ	Indira_Nagar	wm@gmail.com	43
3	h9	4	manchurian	AAK	Garden_Road	h9@gmail.com	50

Time_Taken (In Seconds): 0.9893982410430908

Both the Restaurants Vertical fragments are stored at the same place and they are JOINED before proceeding further. This is an application where JOINS are used.

2. select * from Restaurants where (Rating>3) AND (Rating<5);

a. Query Tree



b. Result

```

->> select * from Restaurants where (Rating>3) AND (Rating<5);

  Name  Rating  Specialty  PK_Custom  Address  Email  Num_Reviews
0  Dominos    4      pizza    AAH    Gachibowli  domi@gmail.com    45
1     h9      4  manchurian  AAK    Garden_Road  h9@gmail.com    50

Time_Taken (In Seconds): 1.2753527164459229
  
```

An application of WHERE clause

3. select count(Name),sum(Num_Reviews),Rating from Restaurants Group By Rating;
- a. Query Tree
The file thirdQuery.md contains the query graph
 - b. Result

```
->> select count(Name),sum(Num_Reviews),Rating from Restaurants Group By Rating;
```

	COUNT_Name	SUM_Num_Reviews	Rating
0	2	95.0	4
1	1	38.0	3
2	1	43.0	5

Time_Taken (In Seconds): 2.1644561290740967

Application of GROUP BY clause

4. select count(Name),sum(Num_Reviews),Rating from Restaurants Group By Rating Having (count(Name)>0) AND (sum(Num_Reviews)>40);
- a. Query Tree
The file fourthQuery.md contains the query graph
 - b. Result

```
->> select count(Name),sum(Num_Reviews),Rating from Restaurants Group By Rating Having (count(Name)>0) AND (sum(Num_Reviews)>40);
```

	COUNT_Name	SUM_Num_Reviews	Rating
0	2	95.0	4
1	1	43.0	5

Time_Taken (In Seconds): 2.227605104446411

Application of HAVING along with GROUP BY

5. select count(Name),sum(Num_Reviews),Rating from Restaurants where (Num_Reviews>43) Group By Rating Having (count(Name)>0) AND (sum(Num_Reviews)>40);
- a. Query Tree
The file fifthQuery.md contains the query graph
 - b. Result

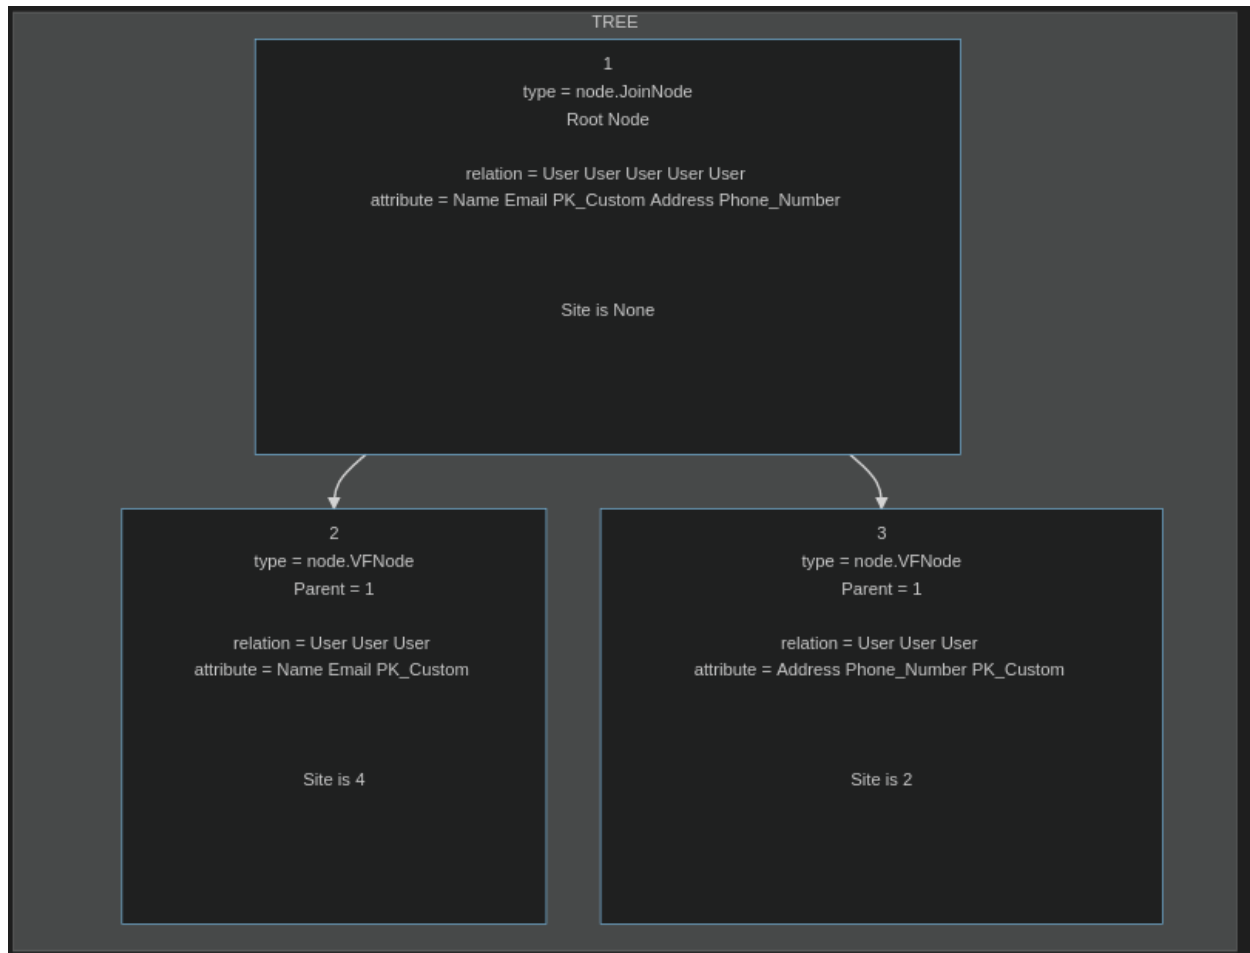
```
->> select count(Name),sum(Num_Reviews),Rating from Restaurants where (Num_Reviews>43) Group By Rating Having (count(Name)>0) AND (sum(Num_Reviews)>40);
```

	COUNT_Name	SUM_Num_Reviews	Rating
0	2	95.0	4

Time_Taken (In Seconds): 3.6412532329559326

Application WHERE along with GROUP BY and HAVING clause

6. Select * from User;
a. Query Tree



- b. Result

```

->> select * from user;

  Name      Email PK_Custom  Address Phone_Number
0   Kartik  krtikgrg@gmail.com    AAA      Sunam  9478077895
1  Aaradhya  aardg@gmail.com      AAB      Delhi  9971352631
2  Priyansh  p@gmail.com          AAC      Delhi  9876543210
3   Harshit  h@gmail.com          AAD  Amritsar  1234567890
4   Aaditya  a@gmail.com          AAE      Jaipur  4561237890
5  Shreyash  s@gmail.com          AAF      Delhi  6543219870
6    Ayush  ag@gmail.com          AAG      Jaipur  7896541230

Time Taken (In Seconds): 3.6980063915252686
  
```

This is an application of inter-site Join, because the fragments for the User table is stored in two separate sites.

Directory Structure

changes.md => File mentioning the changes that have been made since Phase1

complete.md => File representing the final query tree after localization

config.py => File maintaining some global variables which in some sense form the configuration of the system.

data.txt => data that has been added to the sql tables at each site

documentation.md => This file mentions the assumptions that were made while writing the code and also contains some other little details of the code.

executor.py => This file contains the executor class which does nothing but just iterates over the query tree using **DFS** and executes each node.

general_query_syntax.txt => It is a common file that we maintain to help us in coding basically to write notes or points whichever are required.

initial_query_terminal.png => This is a screenshot which shows some commands which were run in succession in the same terminal session.

input.py => File responsible for managing the inputs which we receive in the code.

logs.txt => File which maintains the record of all the functions which were called during the execution of a query. It is a common file for all the queries which are run in a single session.

node.py => File containing the parent class node and all the inherited classes which make it easier to construct and execute the query in form a nice tree structure.

optimization.py => This file contains the optimizer class, which contains nothing but all the functions which are optimizing my query execution.

Post_query_terminal.png => A screenshot in succession to that of **initial_query_terminal.png**

preprocess.py => File containing the functions which will read data from both application and fragmentation schema. It will further process the fragmentation schema to our chosen form. The file also contains other functions now which help us to retrieve the number of tuples in each relation along with each column's size for each relation. There are a couple of more functions which are present in this file.

query.py => File processing the functions of a particular query, all the functions to push down selects and projects including the execution (which will in turn call the executor class method) are present in this file.

server.py => File containing the main loop.

trace.py => File responsible for writing the logs.txt

utility.py => File containing the functions to transfer a relation from one site to another site

report.pdf => File which contains the explanation of the code and the directory structure